

world_model

April 21, 2025

1 World Model

1.1 Imports

```
[2]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor # Good non-linear choice
# Or use MLPRegressor for a neural network:
# from sklearn.neural_network import MLPRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error
from sklearn.preprocessing import StandardScaler
import joblib # To save the trained model
import os
```

1.2 Load Dataset

```
[6]: def load_data(filepath="../dataset/dataset_limited.txt"):
    """Loads the dataset using pandas."""
    try:
        # Assuming the first row is the header
        df = pd.read_csv(filepath)
        print(f"Dataset loaded successfully. Shape: {df.shape}")
        # Optional: Drop rows with NaN values if any occur
        df = df.dropna()
        print(f"Shape after dropping NaNs: {df.shape}")
        return df
    except FileNotFoundError:
        print(f"Error: Dataset file not found at {filepath}")
        return None
    except Exception as e:
        print(f"Error loading dataset: {e}")
        return None
```

```
[ ]: # 1. Load the dataset
dataframe = load_data()
```

Dataset loaded successfully. Shape: (768, 14)

Shape after dropping NaNs: (768, 14)

```
[8]: dataframe.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   distance_red_init      768 non-null    float64
1   angle_red_init         768 non-null    float64
2   distance_green_init    768 non-null    float64
3   angle_green_init       768 non-null    float64
4   distance_blue_init     768 non-null    float64
5   angle_blue_init        768 non-null    float64
6   rSpeed                 768 non-null    int64
7   lSpeed                 768 non-null    int64
8   distance_red_final     768 non-null    float64
9   angle_red_final        768 non-null    float64
10  distance_green_final    768 non-null    float64
11  angle_green_final       768 non-null    float64
12  distance_blue_final     768 non-null    float64
13  angle_blue_final        768 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 84.1 KB
```

1.3 Preprocess Dataset

```
[9]: def prepare_data(df):
      """Separates features (X) and target variables (Y)."""
      # Input Features: initial state (6) + action (2) = 8 features
      X = df.iloc[:, :8].values
      # Target Variables: final state (6) = 6 features
      Y = df.iloc[:, 8:].values
      print(f"Features (X) shape: {X.shape}")
      print(f"Targets (Y) shape: {Y.shape}")
      return X, Y
```

```
[10]: def split_data(X, Y, test_size=0.2, random_state=42):
       """Splits data into training and testing sets."""
       X_train, X_test, Y_train, Y_test = train_test_split(
           X, Y, test_size=test_size, random_state=random_state
       )
       print(f"Training set size: {X_train.shape[0]} samples")
       print(f"Testing set size: {X_test.shape[0]} samples")
       return X_train, X_test, Y_train, Y_test
```

```
[11]: # 2. Prepare Data
      X, Y = prepare_data(dataframe)
```

Features (X) shape: (768, 8)

Targets (Y) shape: (768, 6)

```
[12]: def scale_features(X_train, X_test):  
    """Scales input features using StandardScaler."""  
    scaler = StandardScaler()  
    # Fit scaler ONLY on training data  
    X_train_scaled = scaler.fit_transform(X_train)  
    # Transform both train and test data  
    X_test_scaled = scaler.transform(X_test)  
    print("Features scaled.")  
    return X_train_scaled, X_test_scaled, scaler # Return scaler to save it
```

```
[13]: # 3. Split Data  
X_train, X_test, Y_train, Y_test = split_data(X, Y)  
  
# 4. Scale Features (Important!)  
X_train_scaled, X_test_scaled, scaler = scale_features(X_train, X_test)
```

Training set size: 614 samples

Testing set size: 154 samples

Features scaled.

1.4 Train model

```
[14]: def train_model(X_train, Y_train):  
    """Trains a regression model."""  
    print("Training RandomForestRegressor model...")  
    # Example using RandomForestRegressor  
    # n_estimators: number of trees; random_state: reproducibility; n_jobs=-1: ↵  
    ↵ use all cores  
    model = RandomForestRegressor(n_estimators=100, random_state=42, n_jobs=-1, ↵  
    ↵ oob_score=True)  
  
    # Example using MLPRegressor (Neural Network)  
    # hidden_layer_sizes: tuple defining network structure (e.g., 2 hidden ↵  
    ↵ layers)  
    # max_iter: maximum number of training iterations  
    # alpha: L2 regularization term  
    # learning_rate_init: initial learning rate  
    # early_stopping: stop training if validation score doesn't improve  
    # verbose: print progress  
    # model = MLPRegressor(hidden_layer_sizes=(64, 32), activation='relu', ↵  
    ↵ solver='adam',  
    #                                max_iter=500, random_state=42, early_stopping=True,  
    #                                learning_rate_init=0.001, verbose=True)  
  
    model.fit(X_train, Y_train)
```

```

print("Model training complete.")
# For RandomForest, you can check the Out-of-Bag score as a quick estimate
↳ of performance
if isinstance(model, RandomForestRegressor) and model.oob_score_:
    print(f"Model OOB score: {model.oob_score_:.4f}")
return model

```

```

[15]: # 5. Train Model (using scaled data)
world_model = train_model(X_train_scaled, Y_train)

```

```

Training RandomForestRegressor model...
Model training complete.
Model OOB score: 0.9579

```

```

[16]: def evaluate_model(model, X_test, Y_test):
        """Evaluates the model using MAE and MSE."""
        Y_pred = model.predict(X_test)

        mae = mean_absolute_error(Y_test, Y_pred)
        mse = mean_squared_error(Y_test, Y_pred)
        rmse = np.sqrt(mse) # Root Mean Squared Error

        print("\n--- Model Evaluation ---")
        print(f"Mean Absolute Error (MAE): {mae:.4f}")
        print(f"Mean Squared Error (MSE): {mse:.4f}")
        print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")

        # Optional: Print metrics per output feature
        print("\nMAE per output feature:")
        output_features = [
            'dist_red_final', 'angle_red_final', 'dist_green_final',
            'angle_green_final', 'dist_blue_final', 'angle_blue_final'
        ]
        for i, name in enumerate(output_features):
            mae_feature = mean_absolute_error(Y_test[:, i], Y_pred[:, i])
            print(f" {name}: {mae_feature:.4f}")

        return mae, mse

```

1.5 Evaluate

```

[17]: # 6. Evaluate Model (using scaled test data)
evaluate_model(world_model, X_test_scaled, Y_test)

```

```

--- Model Evaluation ---
Mean Absolute Error (MAE): 22.3253
Mean Squared Error (MSE): 1497.2351

```

Root Mean Squared Error (RMSE): 38.6941

MAE per output feature:

dist_red_final: 39.9570
angle_red_final: 3.9896
dist_green_final: 38.2818
angle_green_final: 7.6712
dist_blue_final: 40.9996
angle_blue_final: 3.0525

[17]: (22.325294401606275, 1497.2351473844744)

1.6 Save model

```
[22]: def save_model_and_scaler(model, scaler, model_filename="world_model.joblib",
    ↪ scaler_filename="scaler.joblib"):
    """Saves the trained model and scaler to disk."""
    try:
        # Ensure the directory exists
        model_dir = "../src/models"
        os.makedirs(model_dir, exist_ok=True)

        model_path = os.path.join(model_dir, model_filename)
        scaler_path = os.path.join(model_dir, scaler_filename)

        joblib.dump(model, model_path)
        joblib.dump(scaler, scaler_path)
        print(f"Model saved to {model_path}")
        print(f"Scaler saved to {scaler_path}")
    except Exception as e:
        print(f"Error saving model/scaler: {e}")
```

```
[23]: # 7. Save Model and Scaler
save_model_and_scaler(world_model, scaler)
```

Model saved to ../src/models/world_model.joblib

Scaler saved to ../src/models/scaler.joblib

```
[24]: # Example prediction (how you'd use it later)
print("\n--- Example Prediction ---")
# Take the first sample from the original test set
sample_X = X_test[0].reshape(1, -1)
sample_Y_actual = Y_test[0]
# Scale the sample using the *saved* scaler
sample_X_scaled = scaler.transform(sample_X)
# Predict using the trained model
sample_Y_pred = world_model.predict(sample_X_scaled)
```

```
print(f"Input State + Action: {sample_X[0]}")
print(f"Actual Final State: {sample_Y_actual}")
print(f"Predicted Final State:{sample_Y_pred[0]}")
```

--- Example Prediction ---

Input State + Action: [215.59621861 -116.56767775 1378.20905064 -172.5025863
1075.89103207

72.16203157 30. 2.]

Actual Final State: [156.33665338 -97.02055051 1296.9236501 -173.7067791
1137.12597569

69.03333994]

Predicted Final State:[227.18284769 -88.00610187 1266.99770549 -161.55511559
1065.22886657

69.61030654]