

world_model_v5

May 6, 2025

1 World Model

1.1 Imports

```
[1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_absolute_error, mean_squared_error
from sklearn.preprocessing import StandardScaler
import joblib
import os
```

```
[2]: # reproducible results
np.random.seed(42)
```

1.2 Load Dataset

```
[3]: def load_data(filepath="../dataset/dataset_v3.txt"):
    """Loads the dataset using pandas."""
    try:
        df = pd.read_csv(filepath)
        print(f"Dataset loaded successfully. Shape: {df.shape}")
        df = df.dropna()
        print(f"Shape after dropping NaNs: {df.shape}")
        return df
    except FileNotFoundError:
        print(f"Error: Dataset file not found at {filepath}")
        return None
    except Exception as e:
        print(f"Error loading dataset: {e}")
        return None
```

```
[4]: # 1. Load the dataset
dataframe = load_data()
```

Dataset loaded successfully. Shape: (3726, 14)
Shape after dropping NaNs: (3726, 14)

```
[5]: dataframe.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3726 entries, 0 to 3725
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   distance_red_init      3726 non-null   float64
1   angle_red_init         3726 non-null   float64
2   distance_green_init    3726 non-null   float64
3   angle_green_init       3726 non-null   float64
4   distance_blue_init     3726 non-null   float64
5   angle_blue_init        3726 non-null   float64
6   rSpeed                 3726 non-null   int64
7   lSpeed                 3726 non-null   int64
8   distance_red_final     3726 non-null   float64
9   angle_red_final        3726 non-null   float64
10  distance_green_final    3726 non-null   float64
11  angle_green_final       3726 non-null   float64
12  distance_blue_final     3726 non-null   float64
13  angle_blue_final        3726 non-null   float64
dtypes: float64(12), int64(2)
memory usage: 407.7 KB
```

1.3 Preprocess Dataset

```
[6]: def prepare_data(df):
      """Separates features (X) and target variables (Y)."""
      # Input Features: initial state (6) + action (2) = 8 features
      X = df.iloc[:, :8].values
      # Target Variables: final state (6) = 6 features
      Y = df.iloc[:, 8:].values
      print(f"Features (X) shape: {X.shape}")
      print(f"Targets (Y) shape: {Y.shape}")
      return X, Y
```

```
[7]: def split_data(X, Y, test_size=0.2, random_state=42):
      """Splits data into training and testing sets."""
      X_train, X_test, Y_train, Y_test = train_test_split(
          X, Y, test_size=test_size, random_state=random_state
      )
      print(f"Training set size: {X_train.shape[0]} samples")
      print(f"Testing set size: {X_test.shape[0]} samples")
      return X_train, X_test, Y_train, Y_test
```

```
[8]: # 2. Prepare Data
      X, Y = prepare_data(dataframe)
```

Features (X) shape: (3726, 8)
Targets (Y) shape: (3726, 6)

```
[9]: def scale_features(X_train, X_test):  
    """Scales input features using StandardScaler."""  
  
    scaler = StandardScaler()  
  
    # Fit scaler ONLY on training data  
    X_train_scaled = scaler.fit_transform(X_train)  
  
    # Transform both train and test data  
    X_test_scaled = scaler.transform(X_test)  
  
    print("Features scaled.")  
  
    return X_train_scaled, X_test_scaled, scaler # Return scaler to save it
```

```
[10]: # 3. Split Data  
X_train, X_test, Y_train, Y_test = split_data(X, Y)  
  
# 4. Scale Features (Important!)  
X_train_scaled, X_test_scaled, scaler = scale_features(X_train, X_test)
```

Training set size: 2980 samples
Testing set size: 746 samples
Features scaled.

1.4 Train model

```
[11]: def train_ridge_regression(X_train, Y_train):  
    """Trains a Ridge Regression model with hyperparameter optimization using  
    ↪GridSearchCV."""  
    print("Training Ridge Regression model with GridSearchCV...")  
  
    # Define the parameter grid for alpha (regularization strength)  
    param_grid = {  
        'alpha': [0.001, 0.01, 0.1, 1.0, 10.0, 100.0, 1000.0], # Wider range ↪  
        ↪of regularization strengths  
        'solver': ['auto', 'svd', 'cholesky', 'lsqr', 'sparse_cg', 'sag', ↪  
        ↪'saga'] # Different solvers for optimization  
    }  
  
    # Create the Ridge regression model  
    ridge = Ridge()  
  
    # Set up GridSearchCV  
    grid_search = GridSearchCV(  

```

```

    estimator=ridge,
    param_grid=param_grid,
    scoring='neg_mean_squared_error', # Optimize for lower MSE
    cv=5, # 5-fold cross-validation
    n_jobs=-1, # Use all available CPU cores
    verbose=1
)

# Fit GridSearchCV on the training data
grid_search.fit(X_train, Y_train)

# Print the best parameters and corresponding score
print("\nGridSearchCV Complete.")
print(f"Best parameters found: {grid_search.best_params_}")
print(f"Best cross-validation score (negative MSE): {grid_search.
↪best_score_:.4f}")

# Return the best model found by GridSearchCV
return grid_search.best_estimator_

# Example usage:
world_model = train_ridge_regression(X_train_scaled, Y_train)

```

Training Ridge Regression model with GridSearchCV...

Fitting 5 folds for each of 49 candidates, totalling 245 fits

GridSearchCV Complete.

Best parameters found: {'alpha': 0.01, 'solver': 'sag'}

Best cross-validation score (negative MSE): -1470.2011

1.5 Evaluate

```

[12]: def evaluate_model(model, X_test, Y_test):
    """Evaluates the model using MAE and MSE."""
    Y_pred = model.predict(X_test)

    mae = mean_absolute_error(Y_test, Y_pred)
    mse = mean_squared_error(Y_test, Y_pred)
    rmse = np.sqrt(mse) # Root Mean Squared Error

    print("\n--- Model Evaluation ---")
    print(f"Mean Absolute Error (MAE): {mae:.4f}")
    print(f"Mean Squared Error (MSE): {mse:.4f}")
    print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")

    # Optional: Print metrics per output feature
    print("\nMAE per output feature:")

```

```

output_features = [
    'dist_red_final', 'angle_red_final', 'dist_green_final',
    'angle_green_final', 'dist_blue_final', 'angle_blue_final'
]
for i, name in enumerate(output_features):
    mae_feature = mean_absolute_error(Y_test[:, i], Y_pred[:, i])
    print(f" {name}: {mae_feature:.4f}")

return mae, mse

```

```

[13]: # 6. Evaluate the Best Model found by GridSearch (using scaled test data) new
evaluate_model(world_model, X_test_scaled, Y_test)

```

```

--- Model Evaluation ---
Mean Absolute Error (MAE): 20.7319
Mean Squared Error (MSE): 1460.9190
Root Mean Squared Error (RMSE): 38.2220

```

```

MAE per output feature:
dist_red_final: 34.2130
angle_red_final: 8.0814
dist_green_final: 33.0882
angle_green_final: 7.7908
dist_blue_final: 34.2454
angle_blue_final: 6.9726

```

```

[13]: (20.73190339516069, 1460.9189664505745)

```

1.6 Save model

```

[14]: def save_model_and_scaler(model, scaler, model_filename="world_model_v5.
↪joblib", scaler_filename="scaler_v5.joblib"):
    """Saves the trained model and scaler to disk."""
    try:
        # Ensure the directory exists
        model_dir = "../src/models"
        os.makedirs(model_dir, exist_ok=True)

        model_path = os.path.join(model_dir, model_filename)
        scaler_path = os.path.join(model_dir, scaler_filename)

        joblib.dump(model, model_path)
        joblib.dump(scaler, scaler_path)
        print(f"Model saved to {model_path}")
        print(f"Scaler saved to {scaler_path}")
    except Exception as e:
        print(f"Error saving model/scaler: {e}")

```

```

[15]: # 7. Save Model and Scaler
      # save_model_and_scaler(world_model, scaler)

[16]: # Example prediction (how you'd use it later)
      print("\n--- Example Prediction ---")

      # Take the first sample from the original test set
      sample_X = X_test[0].reshape(1, -1)
      sample_Y_actual = Y_test[0]

      # Scale the sample using the *saved* scaler
      sample_X_scaled = scaler.transform(sample_X)

      # Predict using the trained model
      sample_Y_pred = world_model.predict(sample_X_scaled)

      print(f"Input State + Action: {sample_X[0]}")
      print(f"Actual Final State:   {sample_Y_actual}")
      print(f"Predicted Final State: {sample_Y_pred[0]}")

```

```

--- Example Prediction ---
Input State + Action: [680.51968108  89.80900352 315.51674792  90.13283211
632.90729297
   4.12401969  22.         -15.         ]
Actual Final State:   [681.67413708  90.6233093  316.80506478  91.88406189
623.33122452
   4.29052584]
Predicted Final State: [683.13227314  85.61458144 319.99822718  86.15491632
631.71629113
   5.66017759]

```